

MouseClaimBench: Integrity-Aware, Non-Compensatory Authorization of Mouse-Brain Claims

Alberto Fernández-Isabel^{a,*}, Isaac Martín de Diego^a, Rubén
Fuentes-Fernández^b

^a*Data Science Laboratory, Research Centre for Intelligent Information Technologies, Rey
Juan Carlos University, C/ Tulipán, s/n, Móstoles, 28933, Madrid, Spain*

^b*Institute of Knowledge Technology, Department of Software Engineering and Artificial
Intelligence, Complutense University of Madrid, C/ Profesor José García Santesmases,
9, Madrid, 28040, Madrid, Spain*

Abstract

Mouse-brain studies combine evidence with distinct meanings. Reproducible execution or accurate prediction may support a bounded result while leaving mechanism, causality, or digital-twin scope unauthorized. MouseClaimBench is an integrity-aware knowledge-based system whose non-compensatory contracts separate domain evidence, structural conformance, provenance integrity, and layer-resolved deficits. Profile v2 links 10 claims, 22 predicate contracts, 60 claim–evidence relations, and 124 observation slots to sources and relation-specific rationales. The profile remains author-defined rather than independently content-validated. Python and a separately implemented Answer Set Programming (ASP) path agree on domain authorization and deficits across 5,677 contract cases. External Shapes Constraint Language (SHACL) independently recovers structural confor-

*Corresponding author

Email address: `alberto.fernandez.isabel@urjc.es` (Alberto Fernández-Isabel)

mance and deficits, while all eight gate combinations satisfy $S \wedge A \wedge I^{\text{clean}}$. A further 55,031 property checks show no violation. Counterfactual tests confirm 9,839 sufficient repairs and 51,480 individual-necessity checks. The integrity gate defines 13 deficit types and blocks 2,550 attacked packages spanning all 255 non-empty compositions of eight historical attack families across 10 claims. Twenty internally coherent forgery controls intentionally escape because no external trust anchor is available. Two pre-access frozen Distributed Archives for Neurophysiology Data Integration (DANDI) protocols preserve one availability refusal and one bounded predictive authorization across 32 mice. A dependence-aware Machine Intelligence from Cortical Networks (MICRONS) application confirms only a local observational structure–function association in two hold-outs. These results establish traceable author-policy execution, not biological truth, criterion validity, human utility, expert consensus, or a complete digital mouse brain.

Keywords: Knowledge-Based Systems, Scientific Claim Authorization, Evidence Provenance, Graph Validation, Mouse-Brain Models, Reproducibility

1. Introduction

Computational neuroscience increasingly depends on reusable datasets, predictive models, connectomic reconstructions, simulation software, and derived statistical artifacts. These resources make analyses more powerful, but
 5 they also create a difficult knowledge-engineering problem. The evidence that supports one statement is frequently insufficient for another. Held-out response prediction does not identify the biological topology that generated

the response. Repeated execution does not provide independent biological replication. A local structure–function association does not establish causal
10 direction or whole-brain coverage. The technical problem addressed in this work is how to encode these distinctions as an executable authorization policy and how to verify that the resulting system does not silently promote a stronger claim.

Mouse-brain modelling is a demanding application because its public re-
15 sources cover distinct biological scales. Allen tracer experiments describe mesoscale connectivity, population recordings characterize behaviour and neural response, Sensorium benchmarks measure predictive performance, and Machine Intelligence from Cortical Networks (MICRONS) links anatomy and function inside a restricted cortical volume [1–4]. The Virtual Mouse
20 Brain, Brain Modeling ToolKit (BMTK), and detailed primary visual cortex (V1) models already provide established simulation capabilities [5–7]. Digital Twin Brain proposals extend this landscape through assimilation and multiscale modelling [8, 9]. None of these resources should be reduced to a common quality score, and a validation layer should not reimplement their
25 simulators.

Existing scientific claim verification primarily retrieves documentary evidence and predicts whether a passage, table, or figure supports a proposition [10–13]. Recent benchmarks also evaluate whether language models can extract and link claim–evidence pairs from complete articles [14]. As-
30 surance cases and knowledge-based systems provide explicit claims, supporting evidence, argument structures, and inference procedures [15–17]. Provenance and graph validation are already standardized through PROV Ontol-

ogy (PROV-O) and Shapes Constraint Language (SHACL) [18, 19]. These components are necessary but do not by themselves answer whether a package containing successful prediction, incomplete lineage, overlapping cohorts, or observational topology evidence authorizes a particular scientific interpretation.

The central design decision of MouseClaimBench is to separate structural conformance, package integrity, and domain authorization. Structural conformance asks whether the required fields and relations exist. Package integrity asks whether profile identity, content hashes, derivation references, independence declarations, and attestations are internally coherent. Domain authorization asks whether every mandatory evidence block for a bounded claim has passed its declared predicate. A package is authorized only when all three layers pass. Missing or failed requirements are returned together, so a topology deficit cannot be hidden by predictive performance and an invalid lineage cannot be hidden by an otherwise complete schema.

This paper develops the integrity-aware authorization version of MouseClaimBench. Its central contribution is an authorization contract that binds four elements that are often evaluated separately: domain evidence semantics, structural admissibility, provenance integrity, and category-complete non-compensatory deficits. The output is deliberately limited to *authorized under profile Π* or *not authorized under profile Π* , accompanied by every violated deficit category and an actionable witness for each detected category. This is not a probability that a scientific statement is true. It is a reproducible decision under one named policy whose content remains open to revision and external criticism.

Research contributions. The paper makes the following contributions:

- A machine-readable knowledge-acquisition record for 10 bounded claims,
60 22 predicate contracts, and all 60 claim–evidence relations. Every relation has a stable identifier, claim-specific necessity rationale, scope, exceptions, sources, and explicit author-policy status.
- An integrity-aware authorization contract that jointly binds evidence-schema admissibility, domain claim requirements, provenance relations,
65 and non-compensatory decisions for computational mouse-brain studies.
- An interoperable Resource Description Framework (RDF)/PROV-O and JSON-LD representation with independently executed SHACL shapes. Python and a separately implemented Answer Set Programming (ASP) path agree on domain decisions and deficits across all 5,677 cases.
70 SHACL independently recovers structural conformance and structural deficits without treating a scientific failure as malformed data.
- A one-edge structural policy perturbation over 221 profiles and 3,094 fixed profile–case evaluations, together with counterfactual tests of explanation sufficiency, minimality, witness coverage, and information retention.
- 75 • An explicit $S \wedge A \wedge I^{\text{clean}}$ final gate with layer-resolved traces, evaluated over all eight logical states, component ablations, a 370-case benchmark and 2,550 attacked packages spanning all 255 compositions of eight historical attack families, a 100-package nine-family regression suite, six direct API edge cases, and 20 coherent-forgery controls.
- 80 • Two pre-access frozen applications to external Distributed Archives for

Neurophysiology Data Integration (DANDI)/Neurodata Without Borders (NWB) resources, a post-outcome threshold-sensitivity analysis that does not replace their criteria, and a dependence-aware MICRONS application. Positive and negative outcomes are preserved without changing endpoints,
85 thresholds, or selection rules after numerical access.

The claimed novelty is the evaluated composition of these elements for artifact-grounded scientific authorization. The work does not claim to invent assurance cases, graph validation, provenance ontologies, formal rule checking, or scientific fact verification. It also does not claim a consensus
90 taxonomy, automatic peer review, biological truth, causal identification from observational data, or a complete digital mouse brain. Section 2 positions the contribution, Section 3 defines the knowledge-based method, Section 4 reports the experimental evaluation, and Sections 5 and 6 discuss its scope and conclusions.

95 2. Related Work

MouseClaimBench lies between scientific claim verification, assurance-case reasoning, semantic validation, and neurocomputational model assessment. Table 1 summarizes the closest methodological families. The comparison is capability-based. It does not assert that each family has one uniform
100 implementation.

2.1. Relationship to the preceding MouseBrainBench study

MouseClaimBench follows the MouseBrainBench validation study [23], but it does not present the earlier model evaluations as new experiments.

Table 1: Capability comparison with the closest methodological families. “Partial” denotes a capability that is narrower, implicit, or implementation-dependent.

Family	Decision object	Graph validation	Package integrity	Scientific authorization
Scientific claim verification [10, 11]	Documentary support or contradiction	No	Source traceability	Partial
Assurance cases [17, 20]	Structured argument for a system property	Partial	Partial	Context-specific
SHACL and PROV-O [18, 19]	Graph conformance and provenance interchange	Yes	Partial	No
Model V&V and reproducibility [21, 22]	Credibility, execution, and context of use	Partial	Partial	Partial
Neural model benchmarks [2, 3]	Prediction, simulation, or biological analysis	No	Dataset- specific	One endpoint
MouseClaimBench	Profile-relative authorization from scientific artifacts	Yes	Thirteen declared integrity- deficit types	Ten bounded claim types

MouseBrainBench introduced a claim-aware validation framework, a diag-
105 nostic Mechanistic Identifiability Score, and a non-compensatory gate over
prediction, reproducibility, topology, direction, local structure–function as-
sociation, and cost. It evaluated synthetic systems and bounded Allen, Sen-
sorium, International Brain Laboratory (IBL), and MICRONS cases. The
present study reuses five frozen result packages from that work only as ret-
110 spective inputs to a different authorization problem. Table 2 separates
inherited material from the contribution evaluated here.

The new contribution is therefore not another analysis of the same five

Table 2: Relationship to the preceding MouseBrainBench study. Reused result packages are retrospective inputs and are not counted as new external validation. “New here” identifies components whose implementation and evaluation constitute the present contribution.

Component	MouseBrainBench	MouseClaimBench
Primary question	Which bounded interpretations are supported by model-evaluation evidence?	Does a versioned evidence package authorize one bounded claim under structural, integrity, and domain gates?
Claim representation	Validation dimensions, evidence contracts, diagnostic score, and conjunctive gate	10 claim types, 22 predicate contracts, 60 direct requirements, and category-complete deficits
Reused artifacts	Allen, Static and Dynamic Sensorium, IBL, and MICRONS evaluations	Five frozen packages mapped retrospectively to profile v2
Authorization engine	Bounded claim gate for the original benchmark	Observation admissibility, typed statuses, complete witnesses, and non-compensatory authorization
Standards layer	Not a central evaluated contribution	RDF/PROV-O, JSON-LD, generated SHACL, and graph-isomorphism checks
Integrity layer	Reproducibility and artifact controls within source analyses	Thirteen declared integrity-deficit types, including eight historical attack families
Independent execution	Original benchmark implementations and controls	Python plus separately implemented ASP, with external SHACL structural validation
Threat evaluation	Not the primary question	Ablations, all 255 declared attack compositions per claim, and coherent-forgery escapes
External protocols	Retrospective public-resource evaluation	Two additional DANDI protocols frozen before numerical access

resources. It is the versioned knowledge profile, admissibility semantics, complete multi-deficit authorization, external graph conformance, relational package-integrity layer, independent rule execution, controlled attack evaluation, and two pre-access DANDI protocols. The dependence-aware MICRONS reanalysis extends uncertainty treatment within the same volume and remains internal reproduction. This separation is also recorded in the repository overlap disclosure and release manifests.

120 2.2. *Scientific claims, knowledge systems, and assurance*

Scientific claim verification commonly decomposes the task into retrieval, evidence selection, and support or contradiction classification [24]. SciFact introduced expert-written claims grounded in scientific abstracts, while MultiVerS and SciFact-Open extended full-document and open-domain reasoning [10, 11, 25]. SciTab, SciVer, SciTrue, PhyVer, and later multimodal approaches include tables, figures, physical measurements, and richer evidence paths [12, 13, 26–28]. CLAIM-BENCH evaluates the extraction and linkage of claim–evidence pairs by language models [14]. These systems address documentary support. MouseClaimBench instead assumes that evidence artifacts and their source predicates have been produced, then evaluates whether their scientific type and provenance satisfy a declared authorization contract. A retrieved sentence cannot substitute for an intervention, an independent cohort, or a topology control.

Rule-based expert systems separate domain knowledge, facts, inference, and explanation [15, 16]. Selective prediction adds a learned reject option when a prediction should not be issued [29]. MouseClaimBench instead returns all typed contract deficits and does not learn an abstention threshold.

Classical validation distinguishes structural correctness from the validity of acquired domain knowledge [30, 31]. This distinction remains central here.
140 Exhaustive tests and multiple solvers can establish that an implementation follows its profile, but they cannot establish that the profile is a complete or universally accepted account of neuroscience.

Assurance cases organize claims, arguments, assumptions, defeaters, and evidence for a defined context of use [17, 20]. Current work also studies
145 assurance cases as attributed graphs and identifies provenance-related structural differences between human- and machine-generated arguments [32]. Compliance-by-construction argument graphs combine typed arguments, deterministic admissibility constraints, and PROV-aligned ledgers for AI-assisted certification records [33]. Claim-level auditability has similarly been
150 proposed for research agents through persistent provenance and protocolized validation [34]. MouseClaimBench does not use a generative model to draft an assurance case. It operationalizes a domain claim profile, tests category-complete deficits across independent execution paths, measures false authorization under controlled package attacks, and evaluates frozen
155 protocols on external scientific artifacts.

2.3. Semantic validation, provenance, and reproducibility

SHACL validates an RDF data graph against a shapes graph and returns a machine-readable validation report [19]. PROV-O represents entities, activities, agents, and derivation relations for interoperable provenance [18].
160 MouseClaimBench reuses both standards. Profile definitions and evidence packages are exported as RDF/JSON-LD, and an external SHACL engine checks field and relation constraints. Consequently, structural graph valida-

tion and provenance interchange are not presented as novel results.

The remaining problem is semantic and relational. A syntactically valid
165 graph may cite the wrong profile, reuse the same artifact as independent
evidence, declare overlapping cohorts as disjoint, contain incompatible attes-
tations, or authorize a claim whose mandatory scientific block failed. These
conditions cross the boundary between graph shape, package integrity, and
domain policy. MouseClaimBench handles them through an additional non-
170 compensatory integrity gate and returns the complete witnesses for every
detected deficit.

Reproducibility guidance, Findable, Accessible, Interoperable, and
Reusable (FAIR) principles, model cards, datasheets, and research-software
standards improve disclosure and reuse [22, 35–39]. NWB provides a com-
175 mon neurophysiology data language and underlies reuse through DANDI
[40]. MouseClaimBench records exact resource versions, asset identifiers,
official hashes, code revisions, protocol hashes, and bounded interpretations.
It also distinguishes clean computational reproduction, within-resource re-
production, and external study replication. These terms cannot be promoted
180 into one another.

Model V&V and uncertainty-quantification frameworks assess implemen-
tation correctness, empirical adequacy, uncertainty, credibility, and context
of use [21, 41–43]. MouseClaimBench does not replace those analyses. It
consumes their resulting artifacts as typed evidence and prevents a passed
185 validation dimension from authorizing another claim whose required evidence
is absent.

2.4. *Validation boundaries in computational mouse-brain research*

Mouse-brain resources differ in anatomical scale and experimental purpose. Allen tracer data and derived connectomes support mesoscale structural analysis [1, 44]. Sensorium and Dynamic Sensorium evaluate neural response prediction under static and video stimuli [2, 3, 45]. Recent work extends predictive generalization and analyzes population-response geometry in visual-cortex digital twins, but prediction remains distinct from unique mechanistic identification [46, 47]. MICRONS provides detailed anatomy and function in visual cortex and supports local structure–function analysis [4, 48]. It is not a measured synaptic connectome of the complete mouse brain.

The Virtual Mouse Brain, BMTK, and detailed V1 models solve simulation and network-construction problems at different scales [5–7]. Observational causal-direction benchmarks show that direction requires assumptions beyond association [49]. Digital Twin Brain and MouseDTB proposals further motivate explicit distinctions among signal prediction, topology, direction, intervention, entity specificity, anatomical coverage, and operational updating [8, 9]. MouseClaimBench does not compete with these simulators. It evaluates which bounded statements their associated evidence packages may authorize.

The targeted review of representative related work did not identify a system combining domain-specific evidence predicates, profile-relative non-compensation, complete multi-deficit traces, separately implemented ASP and external SHACL conformance, package-dependence attacks, prospective external protocols, and mouse-brain artifact applications. This observation

supports differentiation, not global priority or dominance over the representative systems in Table 1. Each component has substantial prior art, and the contribution must be judged on the integration, executable evaluation, and
215 scientific boundaries reported here.

3. Integrity-Aware Knowledge-Based Method


MouseClaimBench receives a versioned authorization profile, a set of evidence blocks, and a manifest of the artifacts from which those blocks were produced. It returns a bounded decision, the complete set of unmet requirements, and a provenance-linked explanation. Fig. 1 shows the complete path.
220 The three validation layers are intentionally separate because a well-formed graph can contain incoherent provenance, and an internally coherent package can still lack the scientific evidence required by a claim.

3.1. Authorization profile and evidence semantics

225 The knowledge profile is represented as

$$\Pi = \langle C, E, Q, O, V, H \rangle, \quad (1)$$

where C is the claim vocabulary, E is the evidence-block vocabulary, $Q(c) \subseteq E$ contains the mandatory blocks for claim c , $O(e)$ contains the observations required when block e declares a pass, V is the profile version, and H is its SHA-256 content identity. The evaluated profile contains 10 claim types, 22
230 evidence-block types, and 60 claim-to-evidence requirements. Table 3 enumerates the claim types, while Fig. 2 shows every required relation. The strict digital-twin claim directly requires 21 blocks. It does not inherit au-



The same claim, evidence blocks, and manifest are evaluated by every gate

Figure 1: Three-gate MouseClaimBench workflow. The same claim, evidence blocks, and manifest determine the RDF/PROV-O representation and enter structural conformance S , domain authorization A , and integrity analysis I . Only the empty-set predicate I^{clean} enters the Boolean composition $A^* = S \wedge A \wedge I^{\text{clean}}$. A failure at any gate prevents authorization, while the output preserves every violated structural, scientific, and integrity category with an actionable witness.

thorization from the other nine claims and is not a synonym for predictive performance.

235 The profile content follows a recorded acquisition procedure rather than being inferred from the conformance benchmark. The procedure first defines bounded claim wording and prohibited promotions. It then identifies distinct evidence functions from methodological and mouse-brain sources, records scope and exceptions, resolves the internal adversarial audit, and verifies
240 exact executable and bibliographic coverage. The machine-readable basis contains one stable record and one claim-specific necessity rationale for each of the 60 direct relations. It also retains the status *source-grounded author*

Table 3: The 10 exact claim types in profile v2. Required-block counts sum to the 60 claim–evidence relations. Contract cases include independent source, rule, and rationale omissions and sum to 5,677.

Code	Claim type	Blocks	Cases
C1	Bounded predictive performance	3	92
C2	Clean computational reproduction	1	13
C3	Within-resource reproduction	2	46
C4	Independent-study replication	3	90
C5	Topology-specific prediction	5	230
C6	Assumption-conditional direction	4	150
C7	Local observational structure–function association	6	321
C8	Intervention-supported effect	5	228
C9	Directed topology-consistent prediction	10	860
C10	Complete entity-specific mouse-brain digital twin	21	3,647
Total		60	5,677

requirement and *not independently validated* for every relation. The source identifiers justify the methodological function of an evidence block. They are
245 not represented as having prescribed the complete profile.

The separate external-review protocol adapts established content-validity procedures [50, 51]. It requires at least seven independent experts, at least five years of relevant experience, representation of three expertise strata, complete 1–4 ratings, item- and scale-level content-validity indices, binary Fleiss
250 agreement, and resolution of every critical veto. The acceptance thresholds were frozen before recruitment. No eligible rater has completed the protocol, so none of these planned statistics enters the present evidence.




Figure 2: Complete profile-v2 claim–evidence matrix. Columns C1–C10 follow Table 3. A filled cell denotes a mandatory evidence block. The 60 filled cells are direct profile relations. In particular, C10 directly requires 21 blocks and does not inherit authorization from C1–C9.

Predicate ownership is recorded separately from claim authorization. Each of the 22 predicate contracts states whether an upstream source adapter, a release adapter, or the authorization adapter evaluates the sci-
 255 entific or operational rule. The contracts expose 124 required observation slots. Source adapters calculate endpoint-specific statuses and preserve the rule, rationale, observations, and artifact identity. The core authorization engine consumes those attestations, checks their schema and provenance
 260 admissibility, and applies $Q(c)$. It does not impose one universal numerical

threshold across prediction, topology, intervention, anatomical coverage, and reproducibility.

An evidence block is

$$b_e = \langle e, a, x, s, \phi, r \rangle, \quad (2)$$

where $a \in \{P, F, U, N, J\}$ denotes passed, failed, unknown, not targeted, or
 265 requires review, x preserves the original-scale observations, s identifies the
 source artifact, ϕ is the predicate executed by the declared source adapter,
 and r is its rationale. The authorization engine validates the attested predi-
 cate record rather than recomputing every source analysis. It does not nor-
 malize correlation, matched-pair differences, software revisions, interventions,
 270 and anatomical coverage into a shared numerical scale. Let

$$\text{present}(z) = \begin{cases} 1, & z \text{ exists and is non-empty,} \\ 0, & \text{otherwise.} \end{cases} \quad (3)$$

A declared pass is admissible only when the block also contains its manda-
 tory observation fields:

$$\text{adm}_\Pi(b_e) = 1[a = P] \text{present}(s) \text{present}(\phi) \text{present}(r) \prod_{o \in O(e)} \text{present}(x_o). \quad (4)$$

Here, $1[\cdot]$ is the indicator function. Consequently, failed, unknown, not tar-
 geted, and requires-review states are all non-admissible. A declared pass is
 275 also non-admissible when its source, predicate, rationale, or a required ob-
 servation is absent. This rule prevents an evidence source from asserting
passed while omitting the target, split, comparator, inferential unit, uncer-
 tainty method, or other observations required by the profile. The complete

scientific deficit set is

$$\Delta_{\Pi}(c, B) = \{e \in Q(c) : \text{adm}_{\Pi}(b_e) = 0\}, \quad (5)$$

280 where an absent block is materialized as unknown. Core authorization is non-compensatory:

$$A_{\Pi}(c, B) = 1[c \in C]1[\Delta_{\Pi}(c, B) = \emptyset]. \quad (6)$$

The engine reports every member of Δ_{Π} with declared and effective states, source, predicate, rationale, observations, and missing fields. It does not select one prioritized explanation or calculate a paper-level score. Category-
 285 complete reporting means that every violated contract dimension is retained without priority masking and has an actionable witness. It does not require exhaustive enumeration of every redundant witness or every distinct cycle within one violation category. Let $T_R(B)$ replace exactly the blocks in a reported repair set R with admissible passing attestations. A category-
 290 complete deficit explanation is counterfactually sufficient and individually necessary when

$$A_{\Pi}(c, T_{\Delta_{\Pi}(c, B)}(B)) = 1 \quad \wedge \quad \forall e \in \Delta_{\Pi}(c, B), \quad A_{\Pi}(c, T_{\Delta_{\Pi}(c, B) \setminus \{e\}}(B)) = 0. \quad (7)$$

This property establishes fidelity to the non-compensatory profile. It does not establish that a human reviewer finds the explanation useful.

3.2. Standards representation and package integrity

295 The profile and evidence package are exported as RDF using PROV-O. Profiles, claims, evidence blocks, artifacts, activities, and decisions receive

stable identifiers. Artifact derivation uses PROV relations, while profile-specific predicates represent mandatory evidence and authorization status. The same graph can be serialized as Turtle or JSON-LD. This layer supports
 300 exchange and inspection. It does not replace the executable profile.

For each profile, MouseClaimBench generates SHACL shapes that require the package and profile identity, all mandatory block records, a status from the allowed vocabulary, source, rule, rationale, and the observations required when a block declares a pass. A third-party pySHACL processor validates the
 305 resulting graph and returns structural deficits. The shapes do not require **passed**. A complete block with **failed** status is therefore structurally conformant but scientifically non-admissible. This distinction permits $S_{\Pi} = 1$ and $A_{\Pi} = 0$ for the same package. SHACL provides an external structural baseline. It does not authorize a scientific claim or make schema validation
 310 a new contribution. In this release, the structural contract applies to graphs emitted by the MouseClaimBench serializer from typed evidence-package objects. It is not presented as a complete validator for arbitrary hostile third-party RDF graphs, where duplicated or contradictory singleton values would require additional uniqueness constraints.

315 The package-integrity layer evaluates relations that cannot be reduced to the presence of individual fields. Let M contain the profile identity, content-addressed artifact records, derivation edges, block-to-artifact lineage, evidence attestations, and declared independence relations. Its integrity deficit set $I_{\Pi}(M, B)$ covers the 13 conditions in Table 4. Artifact identifiers and
 320 block-lineage declarations must be unique. All artifact references from derivations, block lineage, attestations, independence pairs, and disjoint-cohort

pairs must resolve. Block lineage and attestations must reference supplied evidence blocks. Every supplied block requires lineage and at least one attestation, while each direct attestation must agree with the represented block status. Independence and cohort-disjointness relations must be irreflexive. The taxonomy also protects contract identity, artifact content, derivation order, and declared independence. These assets follow established concerns in secure provenance, assurance cases, and the avoidance of pseudoreplication [17, 52, 53].

The cycle check uses an iterative depth-first traversal. Independence is evaluated from content identity and from the composite identity (`study_id`, `data_generation_id`) when both fields are non-empty. The implementation does not assume that `data_generation_id` is globally unique. Duplicate identifiers remain invalid, but validation groups every declaration instead of collapsing them. Reference closure uses all declarations, cycle adjacency uses the union of their derivation edges, independence evaluates all compatible record pairs, and cohort checks use the union declared for each identifier. Input order therefore cannot mask a second deficit category. These checks establish internal consistency, not source trust or scientific independence. Profile identity is checked at both the structural graph and manifest boundaries as explicit defense in depth.

Because $I_{\Pi}(M, B)$ denotes the set of integrity deficits, its Boolean conformance predicate is defined separately:

$$I_{\Pi}^{\text{clean}}(M, B) = 1[I_{\Pi}(M, B) = \emptyset]. \quad (8)$$

Structural conformance remains an independently executed gate. The graph is deterministically derived from the same claim, evidence blocks, and

Table 4: Thirteen-type integrity model and evaluated coverage. O: 370-case original benchmark. C: 2,550-package compositional stress test. E: 100-package extended regression suite. D: six direct API edge-case regressions. O and C use the same eight historical attack families. The 20 trust-boundary controls are separate. These tests do not estimate attack prevalence or adaptive robustness.

Integrity condition	Protected asset	Evaluation	Violated invariant
Profile-version substitution	Authorization contract	O, C	Identifier, version, and profile hash must match
Duplicate artifact identifier	Package namespace	D	Every artifact identifier must be unique within a package
Duplicate block lineage	Evidence trace	D	Every block must have one consolidated lineage declaration
Artifact-hash tampering	Source content	O, C	Declared and observed content identities must agree
Dangling provenance reference	Derivation graph	O, C, E	Every referenced artifact must resolve
Circular provenance	Derivation graph	O, C	Artifact derivation must be acyclic
Duplicate independent artifact	Independence declaration	O, C, D	Independence pairs must be irreflexive and cannot share content or composite generation identity
Overlapping independent cohorts	Inferential units	O, C, D	Cohort-disjointness pairs must be irreflexive and cannot share subjects
Contradictory attestation	Evidence status	O, C	One block cannot have incompatible source attestations
Missing block lineage	Evidence trace	O, C	Every supplied block must resolve to at least one source artifact
Unknown block reference	Evidence trace	E	Lineage and attestations must reference supplied blocks
Block-attestation mismatch	Evidence status	O, C, E	Every direct attestation must agree with the represented block status
Missing block attestation	Evidence trace	E	Every supplied block must have at least one source attestation

manifest evaluated by the other gates:

$$G_{\Pi}(c, B, M) = \text{RDF}_{\Pi}(c, B, M). \quad (9)$$

Let $S_{\Pi}(G) = 1$ exactly when external SHACL validation reports no structural violation against the shapes generated from Π . The field checks in S_{Π} intentionally overlap with defensive admissibility checks in the Python core, but S_{Π} is not absorbed into A_{Π} . Final authorization requires structural
 350 conformance, a complete scientific contract, and an integrity-clean package:

$$A_{\Pi}^*(c, B, M) = S_{\Pi}(G_{\Pi}(c, B, M))A_{\Pi}(c, B)I_{\Pi}^{\text{clean}}(M, B). \quad (10)$$

No control can compensate for another. A dedicated final authorization object executes this equation and preserves separate structural, domain, and integrity traces. Every reported deficit retains its corresponding witness.

355 3.3. Independent execution, domain-gate property checks, and release control

The authorization semantics are implemented through three execution paths with deliberately different scopes. Source-specific Python adapters execute the scientific predicates and construct attested evidence blocks. The primary Python engine evaluates field presence, observation admissibility,
 360 domain requirements, and detailed scientific explanations. The external SHACL path evaluates only the structural contract encoded in the RDF graph. A separately implemented ASP program receives declared statuses, metadata-presence facts, and observation-presence facts. It derives effective states, authorization, and deficits within Clingo for every generated domain-
 365 contract case. Relational package integrity is evaluated by its own Python component before the three gate results are composed. The paths share the

profile specification but not the same inference code. Agreement between Python and ASP, and conformance of SHACL to structural expectations, provide implementation evidence only within their declared scopes. They do
 370 not validate the shared scientific policy.

The domain-gate property contract contains six profile-relative properties of A_{Π} . Soundness requires every domain-authorized claim to have all required admissible blocks. Completeness requires every package with all required admissible blocks to be domain-authorized. Deficit identity requires
 375 the reported set to equal the set defined in Eq. 5. Monotonic degradation requires replacing a passed block with a failed block to preserve or enlarge the deficit set and never create domain authorization. Domain decisions must also be invariant to evidence input order and irrelevant unprofiled blocks. Outside-profile closure is checked separately by requiring every claim out-
 380 side C to remain outside the authorization boundary. Irrelevant-evidence invariance does not apply to arbitrary modifications of a final evidence package. Every supplied block must still satisfy manifest lineage and attestation requirements, even when it is irrelevant to $Q(c)$. These are formalized properties tested over generated packages, not an exhaustive proof over arbitrary
 385 profiles and not statements about biological truth.

Release artifacts record protocol hashes, profile hashes, source revisions, exact public asset identifiers, and official data hashes. A release is rejected when a required artifact is absent, any source revision ends in **-dirty**, an expected benchmark decision changes, or a broad prohibited claim becomes
 390 true. The release gate therefore checks reproducibility and scope together. It cannot guarantee journal acceptance, profile consensus, or human usefulness

of an explanation.

4. Experimental Evaluation

The evaluation is organized into four evidence classes. Knowledge ac-
395 quisition asks whether the author-defined policy is fully traceable. Spec-
ification verification asks whether the implementation follows that policy
and whether its explanations are counterfactually faithful. Controlled threat
evaluation asks which declared integrity violations are detected and which
coherent forgeries remain outside the trust model. Empirical application asks
400 whether bounded decisions are preserved on existing and pre-access frozen
mouse-brain resources. Performance is reported separately and is not used
as scientific validation. Table 5 defines the evidence unit and maximum
interpretation for every question.

4.1. Experimental design and decision rules

405 The acquisition audit compares the executable profile with its machine-
readable basis and bibliography. Acceptance requires exact coverage of 10
claim definitions, 22 evidence blocks and predicate contracts, and 60 claim-
evidence relations. Every relation must have a unique identifier and claim-
specific necessity rationale. Every predicate must state evaluation owner-
410 ship, decision-rule scope, required observations, sources, and exceptions. All
source identifiers must resolve to the bibliography, and the audit must retain
false values for independent content validity and human validation. The his-
torical 29-item v1 panel packet is excluded. A separate v2 protocol contains
92 prospective review items covering the 10 claims, 22 predicates, and 60 re-
415 lations. It requires at least seven eligible independent experts and predefined

Table 5: Evaluation layers, evidence units, and admissible interpretations. Knowledge acquisition, specification verification, controlled threat evaluation, and empirical application remain separate so that software conformance is not reported as scientific content validation.

ID	Question	Evidence unit	Maximum interpretation
K1	Is the scientific policy traceable?	10 claims, 22 predicates, 60 relations, 124 observation slots, and bibliography resolution	Complete author-policy acquisition record, not independent content validity
V1	Are the three gates and their deficits implemented exactly?	5,677 Python/ASP domain cases, 5,677 SHACL structural cases, five JSON-LD round trips, eight gate states, and 10,000 random domain packages	Profile-relative semantic, structural, interoperability, and compositional conformance
V2	How do one-edge changes and repairs affect fixed decisions?	221 profiles on 14 fixed decisions and 10,000 counterfactual packages	Monotonicity and explanation fidelity relative to v2
T1	Does integrity enforce declared package invariants and expose trust limits?	370-case original benchmark, 2,550 attacked compositional packages, 100-package extended regression suite, six direct API edge cases, ablations, and 20 trust-boundary controls	Resistance to declared invariants with explicit coherent-forgery escapes
P1	Is the implementation usable at larger scale?	Up to 10,000 packages and 5,000 artifacts per package	Descriptive performance on one host
A1	Does the profile preserve bounded mouse-brain interpretations?	Five target claims, five strict twin claims, and dependence-aware MICRONS inference	Artifact-grounded authorization under profile v2
A2	Can a frozen contract retain unseen positive and negative outcomes?	Two versioned DANDI resources with pre-access protocols	Prospective workflow evidence within each resource

item content-validity, scale content-validity, agreement, and critical-veto conditions. It has zero completed independent raters and is not counted as evidence.

The contract benchmark generates eight case families from every profile claim. They contain pristine complete packages, one changed status, one
420 omitted mandatory block, one omitted required observation, independent omission of source, rule, or rationale, and all pairwise combinations of the four non-passing states. For a claim with q_c required blocks and m_c mandatory observation fields across those blocks, the number of cases is

$$N_c = 1 + 8q_c + m_c + 16 \binom{q_c}{2}. \quad (11)$$

The terms represent one pristine package, four status mutations, one block
425 omission, and three metadata omissions per block, one omission per observation field, and 16 status combinations per block pair. Table 3 reports every N_c , which sums to 5,677. Each case has an exact expected domain authorization and deficit set defined from Eqs. 4–6. Python and the separately
430 implemented ASP program evaluate all 5,677 domain cases. SHACL evaluates structural conformance and structural deficits over the same packages. Non-passing statuses remain structurally valid, while block, metadata, and required-observation omissions are structural failures. JSON-LD round-trip equivalence is evaluated through graph isomorphism. A separate truth-table
435 benchmark evaluates all eight combinations of S_Π , A_Π , and integrity conformance against Eq. 10.

Domain-gate property-based verification uses 1,000 deterministic random evidence packages for each of the 10 claims. Every package is evaluated for soundness, completeness, exact deficits, input-order invariance, irrelevant-

440 evidence invariance, and monotonic degradation whenever a passed block exists. Outside-profile closure is evaluated separately. Property-based tests generate additional status and observation mutations during the software suite. The reported 55,031 checks refer to the frozen deterministic experiment over A_Π , not to an exhaustive proof over arbitrary future profiles or complete
 445 package-integrity behaviour.

The one-edge structural policy perturbation holds 14 artifact instances fixed: the bounded target and strict complete-twin claim for each of five retrospective cases and two pre-access DANDI cases. It evaluates the baseline profile, all 60 profiles obtained by removing one required relation, and all 160
 450 profiles obtained by adding one currently unrequired evidence block to one claim. Let $x_j = (c_j, B_j)$ denote one fixed claim and evidence package. For variant Π_k , the flip indicator is

$$F_{jk} = 1[A_{\Pi_k}(c_j, B_j) \neq A_\Pi(c_j, B_j)]. \quad (12)$$

Core authorization is used deliberately to isolate the scientific relation policy while manifests remain pristine and graphs remain structurally conformant.
 455 Removing a relation is a weakening perturbation and can only expand authorization. Adding a relation is a conservative extension and can only contract it when the added block is not admissibly passed. The resulting flips are therefore expected monotonicity probes under one-edge changes. They identify decisive relations in the fixed packages, but do not characterize broad
 460 policy dependence, scientific plausibility, or alternative expert profiles.

Counterfactual explanation fidelity is evaluated on another deterministic set of 1,000 packages per claim. Every non-authorized package is repaired by replacing exactly the reported deficits with complete passing attestations.

The repaired package must authorize with no residual deficit. Each repaired
 465 block is then restored individually to its original defective state. Authoriza-
 tion must again fail with that block as the sole remaining deficit, implement-
 ing Eq. 7. Every witness must retain a non-empty block name, source, rule,
 and rationale. Finally, each of the 60 pristine claim–block relations is de-
 graded independently. A single-reason baseline is evaluated descriptively by
 470 counting the simultaneous deficits it would omit. No claim about human
 usefulness is tested.

The original 370-case integrity benchmark evaluates the eight historical
 attack families identified in Table 4. For each of the 10 claims, it contains
 one pristine package, every single attack, and all 28 pairwise combinations,
 475 yielding 360 attacked packages. The structural gate is excluded to isolate
 package-integrity controls, including manifest-level profile identity. For in-
 tegrity configuration k , $I_{\Pi,k}^{\text{clean}}$ applies Eq. 8 to the controls retained by that
 configuration. The domain-plus-integrity decision is

$$A_{\Pi,k}^I(c, B, M) = A_{\Pi}(c, B)I_{\Pi,k}^{\text{clean}}(M, B). \quad (13)$$

Let $\mathcal{D}_{\text{attack}}$ denote the benchmark cases, with claim c_z , blocks B_z , and mani-
 480 fest M_z for case z . The configuration-specific attack false-authorization rate
 is

$$\text{FAR}_{\text{attack}}^{(k)} = \frac{\sum_{z \in \mathcal{D}_{\text{attack}}} 1[z \text{ attacked}] A_{\Pi,k}^I(c_z, B_z, M_z)}{\sum_{z \in \mathcal{D}_{\text{attack}}} 1[z \text{ attacked}]}. \quad (14)$$

We compare the complete domain-plus-integrity configuration with the core
 profile alone, a hash-only baseline, and eight leave-one-control-out configu-
 rations. Exact trace recovery requires the detected deficit codes to equal
 485 every invariant violated by the applied attacks. Full $S_{\Pi} \wedge A_{\Pi} \wedge I_{\Pi}^{\text{clean}}$ compo-
 sition is evaluated separately by the eight-state truth table. Table 8 therefore

measures integrity configurations rather than the final authorization object. The compositional stress test evaluates all $2^8 - 1 = 255$ non-empty subsets for each claim, producing 2,550 attacked packages. A separate nine-family
490 suite contains 10 pristine and 90 reference- or attestation-level regressions. Six direct API cases cover duplicate identifiers, reflexive relations, and both branches of study-scoped data-generation identity. Twenty trust-boundary controls alter content and hashes coherently or forge consistent independence metadata. Their expected authorization exposes the need for external trust
495 anchors rather than an adaptive security guarantee.

The core authorization-engine microbenchmark is a component test measured on one Apple arm64 host with Python 3.12.13. Pristine authorization batches contain 100, 1,000, and 10,000 packages drawn cyclically from all profile claims. Each batch is repeated five times after 100 warm-up decisions. A
500 separate timed component expands one valid digital-twin package to 25, 100, 1,000, and 5,000 uniquely content-addressed artifacts and evaluates its local integrity gate. Median wall-clock time and all repetitions are stored. Downloads, source-predicate execution, graph serialization, and external SHACL validation are excluded. These are component microbenchmarks on one im-
505 plementation and host, not end-to-end scalability or a portable latency guarantee.

Three artifact-evaluation components complement the controlled benchmarks. The first reuses five result packages assembled during the preceding MouseBrainBench study [23] and asks a different question: whether pro-
510 file v2 authorizes their bounded target claim and strict digital-twin claim. These mappings are retrospective and are not counted as prospective vali-

dation. The second applies dependence-aware inference to MICRONS. Its
 endpoint was fixed after discovery and evaluated in two non-overlapping hold-
 out windows within one tissue volume. Directed pair regression controls log
 soma distance, its square, pre- and postsynaptic degree, and coarse cell type.
 515 Dependence is addressed with directed dyadic cluster covariance and 1,000
 simultaneous node-label Freedman–Lane permutations. These hold-outs provide
 internal reproduction, not independent biological replication. The third
 component evaluates the two external DANDI protocols described below.

520 For the MICRONS permutation test, the full model regresses readout-
 location similarity y_{ij} on the directed connection indicator c_{ij} and the fixed
 control vector z_{ij} :

$$y_{ij} = \beta_0 + \beta_1 c_{ij} + z_{ij}^\top \gamma + \varepsilon_{ij}. \quad (15)$$

The reduced model omits c_{ij} . Its residuals are placed in a directed neuron-
 by-neuron array. The connection indicator is residualized against the same
 525 controls, and the residualized coefficient is the test statistic. Each permuta-
 tion draws one permutation of neuron identifiers and applies it simultaneously
 to sender and receiver indices of the reduced-model residual array. This pre-
 serves reciprocal array structure and the incidence pattern induced by shared
 neurons under relabelling. The discovery window fixes the positive direction
 and the analysis specification. Its one-sided permutation value is reported
 530 descriptively and has no confirmatory interpretation. The dyadic regression
 is reported with a two-sided test and 95% interval in every window. The con-
 firmatory one-sided permutation test is subsequently evaluated in the two
 pre-fixed hold-outs. Its finite-permutation value is $(1 + K)/(1 + R)$, where
 535 K is the number of permuted coefficients at least as large as the observed

positive coefficient and $R = 1000$. Discovery and the two hold-outs use reproducible streams with seeds 2026080201, 2026080202, and 2026080203. A hold-out passes only when $\hat{\beta}_1 > 0$, $p_{\text{dyadic}} \leq 0.05$, and $p_{\text{perm}} \leq 0.05$. Confirmation requires this conjunction in both fixed hold-outs after a positive
540 discovery direction. This procedure assumes exchangeability of the reduced-model residual array under simultaneous node relabelling.

Two external DANDI protocols were committed before numerical asset access. The availability protocol was frozen in commit **b59a497** and the predictive protocol in **d039e66**. Their SHA-256 identities are retained in
545 the result artifact, and the numerical result was committed later. This Git history is temporal evidence for the declared ordering, not third-party preregistration. DANDI:001176 was selected from metadata describing 29 subjects and simultaneous cholinergic axon, cortical acetylcholine (ACh), pupil, and treadmill signals [54]. The protocol required at least 20 usable subjects and
550 prohibited threshold or endpoint repair. After catalog inspection, 28 schema-valid sessions came from only 5 subjects. The predictive model was therefore not executed.

DANDI:000039 contains contrast-tuning recordings associated with a published visual-cortex study [55]. Before inspecting its asset catalog or values,
555 the protocol fixed one lexicographically first **behavior+ophys.nwb** asset per subject, a trial-level population mean event response, Ridge regression with $\alpha = 1$, and an intercept comparator. For contrast q and direction θ in degrees, the exact five regressors are $\log(1 + 100q)$, $\sin(\theta)$, $\cos(\theta)$, $\log(1 + 100q) \sin(\theta)$, and $\log(1 + 100q) \cos(\theta)$. Angular functions are evaluated after conversion
560 to radians. The first 60% of trials train the model, the next 20% remain

reserved, and the final 20% form the test partition. The intermediate partition is deliberately unused. It does not affect model fitting, model selection, threshold definition, or final evaluation. The `StandardScaler` is fitted on the training partition only and then applied unchanged to the held-out partition.

565 Ridge $\alpha = 1$ is fixed and no tuning occurs after asset access. The per-mouse held-out metric is Pearson correlation. Authorization requires at least 25 subjects, median held-out correlation at least 0.10, a positive 95% mouse-bootstrap lower bound, at least 65% positive subjects, and trial-pooled mean squared error below the intercept. For M mice with n_m held-out trials, this
570 metric is

$$\text{MSE}_{\text{pool}} = \frac{\sum_{m=1}^M \sum_{t=1}^{n_m} (y_{mt} - \hat{y}_{mt})^2}{\sum_{m=1}^M n_m}. \quad (16)$$

The implementation stores and compares the two squared-error numerators because their denominator is identical. The resulting metric weights mice in proportion to their number of held-out trials. All conditions must pass. The bootstrap uses protocol seed 2026080202 and resamples mice with replacement
575 5,000 times. Each replicate calculates the median mouse-level Pearson correlation. The reported lower bound is its 2.5th percentile. It therefore targets population stability across mice rather than trial-level independence.

These thresholds are author-defined operational acceptance criteria rather than criterion-validated biological standards. Twenty-five mice was intended
580 to preserve a multi-mouse endpoint and a meaningful mouse-level bootstrap, but it was not obtained from a formal power calculation. The 0.10 median correlation requires a non-trivial predictive signal, the positive lower bound requires population-level sign stability, the 65% rule requires a majority of positive mice while tolerating heterogeneity, and the intercept comparison

585 requires incremental predictive value. The 20-subject availability boundary for DANDI:001176 similarly protects the intended cross-subject design but is not a sufficiency standard. A post-outcome sensitivity analysis changes one threshold at a time while holding the other frozen conditions fixed. It exposes decision boundaries and cannot retrospectively calibrate or replace
590 either protocol.

4.2. Knowledge traceability, one-edge policy perturbation, and explanation fidelity

The acquisition audit exactly covers the 10 claims, 22 evidence blocks, 22 predicate contracts, 60 claim–evidence relations, and 124 required obser-
595 vation slots. All 60 relation identifiers and necessity rationales are unique, and all 20 source identifiers resolve to the bibliography. The ledger separates upstream predicate execution from authorization checks and marks every relation as a source-grounded author requirement without independent consensus. This satisfies traceability of the current policy. It does not establish
600 that the policy is complete or optimal.

The one-edge analysis evaluates 221 profiles and 3,094 profile–case combinations. The baseline authorizes four of the 14 fixed decisions. Among the 60 weakening variants, only removal of the `data_quality` requirement from bounded prediction changes an authorization. It converts the Dynamic Sen-
605 sorium target from a documented quality refusal into an authorization. This isolates the requirement that is decisive for that case rather than showing that it should be removed. Among 160 conservative extensions, 52 variants change at least one decision and produce 71 contractions. These contractions follow the non-compensatory monotonicity of the system because an other-

610 wise authorized case has no admissible attestation for the newly added block. Ninety-six variants change at least one deficit set even when authorization does not change. The counts are implementation context, not independent scientific evidence for the added relations.

Table 6: One-edge structural policy perturbation across 14 fixed artifact decisions. Expansions are refusals authorized after removing one requirement. Contractions are authorizations withdrawn after adding one block. The monotonicity probe identifies decisive relations in these packages. The variants are not scientifically validated candidate profiles.

Policy set	Variants	Runs	With flips		Flips	Interpretation
Baseline profile	1	14	–		0	Four bounded authorizations
Remove one required relation	60	840	1	1 expansion		Removing data quality authorizes the otherwise refused Dynamic Sensorium target
Add one unrequired block	160	2,240	52	71 contractions		Positive cases lose authorization when the added block lacks an admissible attestation
Total	221	3,094	53		72	One-at-a-time structural policy envelope

Counterfactual fidelity is evaluated on 10,000 packages. Of the 9,839
615 non-authorized packages, repairing every reported deficit is sufficient for authorization with no residual deficit. Restoring each deficit individually produces 51,480 minimality checks, all of which refuse authorization with the exact retained deficit. The same 51,480 deficit records retain non-empty witnesses, and all 60 pristine single-block degradations return the expected
620 sole deficit. There are 8,632 packages with simultaneous deficits. A single-reason trace would retain 9,839 deficits and therefore omit the remaining

51,480 − 9,839 = 41,641. This is arithmetic information retention under the generated packages, not independent evidence about explanation quality or human decisions.

625 4.3. Domain-gate conformance and property-based verification results

Table 7 reports the conformance results. Python produces no false authorization or rejection in 5,677 cases and recovers every domain deficit set. The independent ASP implementation matches these outcomes while deriving admissibility inside Clingo. External SHACL returns every expected structural
630 decision and deficit, including status-only failures that remain structurally valid but scientifically unauthorized. All eight combinations of structural, domain, and integrity conformance satisfy Eq. 10. The exported 318-triple RDF profile represents all claims and evidence blocks. JSON-LD round trips preserve graph isomorphism in five packages spanning pristine, scientific-failure,
635 structural-failure, rich-manifest, and digital-twin-style cases.

Python and ASP agree on domain authorization. SHACL agrees with structural expectations rather than domain decisions and recomputes neither source predicates nor relational integrity. Structural validation is therefore a baseline, while the contribution lies in non-compensatory composition, domain
640 requirements, package integrity, controlled attacks, and applications.

The 10,000 deterministic random packages produce 55,031 domain-gate property checks with no violation. Soundness, completeness, deficit identity, input-order invariance, and irrelevant-evidence invariance each pass 10,000 checks. Monotonic degradation passes all 5,030 applicable checks. Outside-
645 profile closure also passes. These results establish the operational semantics of A_{Π} for profile v2. They do not establish final-package invariance or prove

Table 7: Conformance and domain-gate property results. Exactness is relative to the frozen profile-v2 contract, and each engine is reported only for its evaluated comparison set.

Evaluation	Cases/checks	Errors	Result
Python contract mutations	5,677	0	Exact domain authorization and deficits
Independent ASP conformance	5,677	0	100% domain-status and deficit agreement
External SHACL conformance	5,677	0	Exact structural conformance and structural deficits
Three-gate composition	8 states	0	Exact $S \wedge A \wedge I^{\text{clean}}$ truth table
JSON-LD round trips	5 graphs	0	Isomorphic RDF graphs
Domain-gate property benchmark	55,031	0	Six core properties and outside closure hold

that the evidence requirements are scientifically complete.

4.4. Integrity attacks, ablation, and component microbenchmarks

Under the declared deterministic threat model, the complete domain-plus-
650 integrity configuration returns the exact integrity-deficit trace in all 370 packages. It produces no false rejection on the 10 pristine packages and no false authorization on the 360 attacked packages. The core profile does not inspect package integrity and therefore authorizes all 360 attacked packages. Hash checking alone blocks content tampering but authorizes 280 packages
655 affected by profile substitution, broken or cyclic lineage, false independence, cohort overlap, contradictory attestations, or missing block provenance. The

structural gate is excluded from these counts, so manifest-level profile substitution remains visible instead of being rejected earlier by structural identity matching.




Figure 3: False authorizations among 360 attacked packages across 10 profile claims, where lower values are better. The structural gate is excluded to isolate the domain-plus-integrity configurations. The profile-only baseline authorizes all 360 attacked packages. The hash-only baseline authorizes 280 packages involving identity, lineage, dependence, conflict, or provenance. The full integrity configuration blocks every declared attack. Seven original control omissions expose 10 packages each. Removing the contradiction check alone exposes none because block-attestation status consistency detects the same contradictory packages.

660 The leave-one-out result distinguishes unique necessity from defensive overlap. Removing any of seven original controls creates 10 false authorizations, corresponding to the single-attack package for that omitted control in each claim. Removing only the contradictory-attestation check creates none

because the same package also violates block–attestation status consistency.

665 Pairwise attacks containing another enabled control remain blocked. Table 8 gives the exact values, while Fig. 3 shows the contrast between incomplete baselines and the complete domain-plus-integrity configuration. The result demonstrates conditional necessity and one explicit redundancy under the declared attack construction. It does not imply sufficiency for every real

670 provenance threat.

Table 8: False authorizations and false rejections for domain-plus-integrity configurations in the 370-case integrity benchmark. The structural gate is excluded to isolate manifest-level controls. Seven original controls are uniquely necessary for their constructed single-attack families. Removing the contradiction check alone produces no false authorization because block–attestation status consistency independently detects the same contradictory package.

Integrity configuration	False authorizations / 360	False rejections / 10
Profile without integrity	360	0
Hash-only integrity	280	0
Full integrity gate	0	0
Without profile identity	10	0
Without content hash	10	0
Without resolved lineage	10	0
Without cycle detection	10	0
Without duplicate-independence check	10	0
Without cohort-overlap check	10	0
Without contradictory-attestation check	0	0
Without block lineage	10	0

The compositional stress test evaluates 2,560 packages, comprising 10
pristine packages and 2,550 attacked packages. For each claim, the latter
span all 255 non-empty compositions of the eight historical attack families.
The complete domain-plus-integrity configuration produces no false autho-
675 rization and returns every implied invariant deficit. The extended nine-family
regression suite returns exact traces across its 100 packages, with no false au-
thorization among 90 attacked packages and no false rejection among 10
pristine packages. The six separate direct API regressions return structured
refusals for duplicate and reflexive declarations and the expected pass or re-
680 fusul under the composite generation-identity rule. All 20 internally coherent
forgeries remain authorized because no externally verifiable contradiction ex-
ists. The system therefore detects inconsistency with declared invariants, but
it does not authenticate self-consistent source metadata.

Batch throughput ranges from 27,653 to 52,816 authorizations per second.
685 Median time is 0.00189 s for 100 packages, 0.03253 s for 1,000, and 0.36163 s
for 10,000. A valid package requires 0.149 ms with 25 artifacts, 0.240 ms with
100, and 5.328 ms with 5,000. These measurements describe the observed
range on one host and are not an asymptotic complexity result. All pristine
decisions remain authorized at every scale.




Figure 4: MouseClaimBench component microbenchmarks on one Apple arm64 host using medians from five repetitions after warm-up. The left panel times local core authorization for batches of 100, 1,000, and 10,000 pristine packages. The right panel times the local integrity gate for one valid package containing 25, 100, 1,000, or 5,000 artifacts. Downloads, source predicates, graph serialization, and external SHACL execution are excluded. The measurements are not end-to-end throughput.

Table 9: MouseClaimBench component microbenchmarks on one Apple arm64 host. Every value is the median of five repetitions after warm-up. Downloads, source predicates, graph serialization, and external SHACL execution are excluded.

Workload	Size	Median time (ms)	Throughput
Evidence-package batch	100 packages	1.893	52,816 packages/s
Evidence-package batch	1,000 packages	32.529	30,742 packages/s
Evidence-package batch	10,000 packages	361.629	27,653 packages/s
Single digital-twin package	25 artifacts	0.149	Not pooled
Single digital-twin package	100 artifacts	0.240	Not pooled
Single digital-twin package	1,000 artifacts	1.344	Not pooled
Single digital-twin package	5,000 artifacts	5.328	Not pooled

690 4.5. *Mouse-brain artifacts and prospective external applications*

The retrospective mapping evaluates five bounded targets and five strict digital-twin claims. Static Sensorium authorizes prediction, IBL authorizes topology-specific behavioural prediction, and MICRONS authorizes one local observational association. Allen directed topology lacks several required
695 blocks, while Dynamic Sensorium retains a response-reliability review state. No case authorizes the strict digital-twin claim.

The dependence-aware MICRONS analysis changes the status of the local association without changing its scope. Table 10 reports cohort size, pair counts, effect estimates, uncertainty, and both inferential checks. No per-
700 mutated statistic is at least as large as the observed positive statistic in any window. With 1,000 permutations, the finite-permutation corrected value is therefore $p = 1/1001 \approx 0.001$ in each window. The discovery value is descriptive because that window selected the positive direction. Both prefixed hold-outs independently pass the positive coefficient, two-sided dyadic,
705 and one-sided permutation criteria. Shared-neuron dependence is addressed more conservatively than naive pair-level inference. All windows nevertheless come from one MICRONS volume, and readout-location similarity is a derived functional descriptor rather than a causal measurement.

The DANDI outcomes are asymmetric. DANDI:001176 fails availability
710 with 5 usable subjects against a frozen minimum of 20, so no predictive fit is reported. DANDI:000039 yields 32 usable mice, median held-out correlation 0.310, mouse-bootstrap 95% interval [0.207, 0.369], and 93.75% positive subjects. Model and intercept squared-error sums are 11.579 and 12.244 over 6,193 held-out trials, a 5.43% relative reduction. All five conditions pass,

Table 10: Dependence-aware MICRONS inference for the fixed all-pairs/readout-location endpoint. Candidate and connected counts refer to directed pairs. Confidence intervals use the same unit-based degrees of freedom as the directed dyadic test. Discovery fixes the positive direction, so its one-sided permutation value is descriptive. Confirmation requires $\hat{\beta} > 0$, two-sided $p_{\text{dyadic}} \leq 0.05$, and one-sided $p_{\text{perm}} \leq 0.05$ in both pre-fixed hold-outs.

Cohort	Units	Pairs (connected)	$\hat{\beta}$ (SE)	95% CI	p_d	p_p
Discovery	1,000	999,000 (2,095)	0.0106 (0.00200)	[0.00668, 0.01451]	1.39×10^{-7}	0.0010
Hold-out 1	992	983,072 (1,926)	0.0178 (0.00209)	[0.01372, 0.02192]	5.78×10^{-17}	0.0010
Hold-out 2	999	997,002 (1,922)	0.0117 (0.00205)	[0.00763, 0.01569]	1.76×10^{-8}	0.0010

715 authorizing only bounded population-response prediction. A later audit confirms that every held-out contrast, direction, and joint condition occurs in training for every mouse. It changes no split, model, threshold, or decision.

Table 11: Post-outcome, one-at-a-time sensitivity of the author-defined DANDI acceptance criteria. The frozen threshold is the prospective decision rule. The last column reports the nearest tested decision boundary or, when no boundary is crossed, the observed outcome. It is not a recommended replacement or a criterion-valid standard.

Criterion	Frozen	Observed	Sensitivity boundary or outcome
DANDI:000039 usable mice	≥ 25	32	≥ 33
Median mouse correlation	≥ 0.10	0.310	≥ 0.35
Bootstrap lower bound	> 0	0.207	> 0.25
Positive-mouse fraction	≥ 0.65	0.938	≥ 0.95
Relative squared-error gain	> 0	5.43%	$> 7.5\%$
DANDI:001176 usable subjects	≥ 20	5	Negative at every tested minimum ≥ 10 and available only at ≤ 5

One-at-a-time sensitivity preserves the frozen DANDI:000039 decision and exposes its finite margins. The decision fails above 32 required subjects,
720 at median correlation 0.35, bootstrap lower bound 0.25, positive-subject frac-

tion 0.95, or required squared-error improvement 7.5%. DANDI:001176 remains unavailable for every tested minimum of at least 10 and would pass only after the prohibited relaxation to five. These post-outcome checks provide no criterion validity.




Figure 5: Prospective DANDI outcomes under rules fixed before numerical asset access. DANDI:001176 provides 5 usable subjects against the predeclared minimum of 20, so no model is fitted. DANDI:000039 passes every frozen condition across 32 mice, with median held-out correlation 0.310, subject-bootstrap lower bound 0.207, 93.8% positive mice, and relative MSE reduction 0.054. Red diamonds mark the respective thresholds of 0.10, 0, 0.65, and 0. Bar heights are criterion-specific values and are not intended for direct cross-metric comparison. Only bounded population-response prediction is authorized.

725 The positive result is not a state-of-the-art model claim. Ridge regression was chosen because the experiment tests prospective authorization rather than neural architecture. The negative result is equally part of the method evaluation. Together they show that a frozen knowledge contract can preserve both authorization and refusal without post-outcome threshold repair.

730 The sensitivity analysis was executed later and cannot alter either decision. The final release verifies all source revisions, profile and protocol hashes, exact

public DANDI asset identifiers, and the absence of **-dirty** artifacts.

Table 12: Artifact-grounded and prospective mouse-brain decisions. Authorization never extends beyond the claim stated in the second column.

Resource	Bounded target claim	Decision	Decisive result or deficit
Allen VBN	Directed topology and prediction	No	Topology and direction fail, with further missing blocks
Static Sensorium	Bounded predictive performance	Yes	Five-dataset comparison passes declared prediction and quality rules
Dynamic Sensorium	Bounded predictive performance	No	Response reliability requires review
IBL	Topology-specific behavioural prediction	Yes	Locked two-split comparison passes, without exclusive superiority
MICRONS	Local structure-function association	Yes	Direction fixed in discovery. Both pre-fixed hold-outs pass dyadic and node-label permutation criteria
DANDI:001176	Cortical ACh prediction	No	5 usable subjects against a frozen minimum of 20
DANDI:000039	Population contrast-response prediction	Yes	32 mice, median $r = 0.310$, bootstrap lower bound 0.207
Strict digital twin	Entity-specific whole-brain twin	No	Zero authorizations across all evaluated resources

5. Discussion

The results support MouseClaimBench as a bounded knowledge-based au-
735 thorization layer. They do not support a claim that the system determines
scientific truth. The system prevents evidence valid for one claim from being
promoted into another that requires different evidence, using named require-
ments, package integrity, and category-complete deficit traces.

5.1. Scientific interpretation and methodological novelty

740 External SHACL reproduces every structural decision and deficit. Mouse-
ClaimBench therefore claims no novelty for required-field validation, prove-
nance representation, or JSON-LD export. Assurance cases already repre-
sent claims and evidence, and scientific verification retrieves documentary
support. The contribution that remains is their evaluated composition with
745 domain predicates, relational integrity, non-compensatory authorization, and
unmasked deficit categories. Python and the separately implemented ASP
path agree on domain outputs, SHACL permits scientifically failed but well-
formed packages to conform, and the final object executes $S \wedge A \wedge I^{\text{clean}}$ while
retaining each layer’s deficits.

750 The one-edge result is an expected monotonicity probe, not broad policy
validation. Fifty-nine removals are not singly decisive because other deficits
remain, while one removal exposes data quality as the sole Dynamic Senso-
rium barrier. The 71 contractions caused by added unprovided blocks do
not establish that those additions are scientifically appropriate. Similarly,
755 counterfactual repair establishes formal sufficiency, individual necessity, and
witness retention, not human usefulness. The 41,641 deficits omitted by a

one-reason baseline quantify information loss only.

The 2,550 attacked packages cover every composition of the eight historical families for every claim, but remain specification-driven. The nine-family
760 suite and direct API regressions extend this controlled scope. All 20 coherent forgeries escape because synchronized false content, hashes, or independence metadata create no observable contradiction. The result identifies both controlled robustness and the trust boundary. It does not address digital signatures, archive compromise, adaptive attacks, or unmodelled dependence.

765 The prospective applications test protocol discipline rather than independent discovery. DANDI:000039 uses a deliberately simple Ridge model and author-defined criteria without established criterion validity. DANDI:001176 retains its availability refusal when only 5 of 29 listed subjects meet the frozen schema. The post-outcome sensitivity analysis exposes decision margins with-
770 out replacing either rule. In MICRONS, dyadic covariance shows that naive pair-level inference is anti-conservative. Discovery fixes the direction, and both pre-fixed hold-outs confirm the local observational association. One cortical volume cannot establish a synaptic mechanism, external replication, or a complete visual-cortex model.

775 *Lessons learned.* Four lessons follow. Structural conformance should use a standard engine. Provenance must record dependence claims, not only locations. Negative outcomes should remain release artifacts. Category-complete traces should retain every violated dimension and an actionable witness, but need not enumerate redundant occurrences. Their value to human reviewers
780 remains unmeasured.

5.2. Limitations and threats to validity

Construct validity remains the principal limitation. Profile v2 is author-defined and literature-grounded, with inspectable claims, predicates, relations, rationales, exceptions, and sources. No independent panel has established their necessity, sufficiency, or completeness, and the frozen 92-item review protocol has zero eligible responses. Conformance, counterfactual fidelity, and perturbation are therefore profile-relative rather than evidence of scientific consensus.

Internal validity is stronger for software behaviour than scientific content. Expected mutation, counterfactual, and attack outcomes follow from the declared semantics. Multiple engines reduce implementation risk but share the profile specification. One-edge variants cover local monotonic behaviour, property checks sample finite packages, the attack generator is deterministic, and coherent false metadata remains unobservable. The DANDI sensitivity analysis is descriptive and provides no criterion validity.

The applications are heterogeneous rather than biological replications. Five retrospective packages come from MouseBrainBench [23]. The MICRONS windows share one volume, the DANDI studies use different measurements, DANDI:000039 tests bounded prediction rather than model superiority, and DANDI:001176 tests availability. Mouse-level bootstrap and dyadic covariance address declared dependence units but retain exchangeability assumptions. The hold-outs and frozen Git protocols were not registered externally.

External validity is not established outside computational mouse-brain evidence. Loading another profile is architectural extensibility, not empirical

generality. Human comprehensibility and decision utility were not evaluated. Performance is limited to component microbenchmarks on one Apple arm64 host and excludes data access, source predicates, graph serialization, and external SHACL execution.

810 6. Conclusions and Future Work

MouseClaimBench provides an integrity-aware knowledge-based method for authorizing bounded claims from computational mouse-brain evidence. Its contracts bind graph conformance, package integrity, domain semantics, and category-complete non-compensatory deficits. Each decision is tied to a
815 named profile and retains every violated category with an actionable witness.

The evaluation separates acquisition, specification verification, controlled threats, and empirical application. The ledger covers all 60 claim–evidence relations and 22 predicate contracts without claiming external consensus. Python and the separately implemented ASP path agree across 5,677 do-
820 main cases, external SHACL recovers the expected structural outcomes, and the final gate reproduces all eight states of $S \wedge A \wedge I^{\text{clean}}$. One-edge perturbation and counterfactual repair establish implementation conformance and explanation fidelity relative to profile v2, not scientific completeness.

The integrity evaluation blocks all 2,550 attacked packages under the de-
825 clared families, while all 20 coherent forgeries expose the need for external trust anchors. One prospective DANDI protocol retains an availability refusal and another authorizes only population-response prediction across 32 mice. Dependence-aware MICRONS inference confirms one local observational association in two pre-fixed hold-outs. No case authorizes causality,

830 external biological replication, whole-brain coverage, or a complete entity-specific mouse-brain digital twin.

Future work requires independent profile assessment, human evaluation of traces, signed attestations, adaptive attacks, and externally registered endpoints. The current conclusion remains narrower: prediction, reproducible
835 execution, and graph validity cannot compensate for missing scientific evidence or incoherent provenance under the declared profile.

Declaration of Competing Interest

The authors declare no known competing financial interests or personal relationships that could have appeared to influence the work reported in this
840 paper.

CRedit Authorship Contribution Statement

Alberto Fernández-Isabel: Conceptualisation, Methodology, Software, Investigation, Data curation, Writing–original draft, Writing–review and editing.
Isaac Martín de Diego: Validation, Formal analysis, Investigation, Writing–
845 original draft, Writing–review and editing, Project administration.
Rubén Fuentes-Fernández: Methodology, Validation, Writing–original draft, Writing–review and editing.

Data and Code Availability

The source code, knowledge profiles, tests, frozen protocols, lightweight
850 result artifacts, and manuscript are maintained in the public MouseClaimBench repository at <https://github.com/Fiutten/MouseClaimBench>.

The software snapshot used for this revision is released as tag v0.12.4 at <https://github.com/Fiutten/MouseClaimBench/tree/v0.12.4>. Each regenerated revision artifact records the clean source commit used for its execution. Raw Allen, Sensorium, IBL, MICRONS, and DANDI data remain governed by their providers and are not redistributed. The prospective artifact manifests retain public identifiers, versions, sizes, and official hashes. Every reported result is linked to a recorded Git revision, and artifacts marked as dirty are excluded from the release gate. A persistent DOI has not been assigned, so the repository and tagged commit are reported without implying DOI-backed archival permanence.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During preparation of this work, the authors used OpenAI Codex, an Artificial Intelligence (AI)-assisted tool, for software assistance and language editing. The authors reviewed and edited all generated material and take full responsibility for the content of the publication.

Funding

This work was supported by the Spanish MICIU under the PDI program in the XMIDAS project [PID2021-122640OB-I00]. The funder had no role in study design, data collection, analysis or interpretation, manuscript preparation, or the decision to submit the article for publication.

References

- [1] S. W. Oh, J. A. Harris, L. Ng, B. Winslow, N. Cain, S. Mihalas, et al.,
875 A mesoscale connectome of the mouse brain, *Nature* 508 (2014) 207–214.
[doi:10.1038/nature13186](https://doi.org/10.1038/nature13186).
- [2] K. F. Willeke, P. G. Fahey, M. Bashiri, L. Pede, M. F. Burg, C. Blessing,
S. A. Cadena, Z. Ding, K.-K. Lurz, K. Ponder, T. Muhammad, S. S.
Patel, A. S. Ecker, A. S. Tolias, F. H. Sinz, The Sensorium competition
880 on predicting large-scale mouse primary visual cortex activity, *arXiv*
(2022). [arXiv:2206.08666](https://arxiv.org/abs/2206.08666), [doi:10.48550/arXiv.2206.08666](https://doi.org/10.48550/arXiv.2206.08666).
- [3] P. Turishcheva, P. G. Fahey, L. Hansel, R. Froebe, K. Ponder, M. Vys-
trčilová, K. F. Willeke, M. Bashiri, E. Wang, Z. Ding, A. S. Tolias,
F. H. Sinz, A. S. Ecker, The Dynamic Sensorium competition for pre-
885 dicting large-scale mouse visual cortex activity from videos, *arXiv* (2023).
[arXiv:2305.19654](https://arxiv.org/abs/2305.19654), [doi:10.48550/arXiv.2305.19654](https://doi.org/10.48550/arXiv.2305.19654).
- [4] The MICrONS Consortium, Functional connectomics spanning multiple
areas of mouse visual cortex, *Nature* 640 (2025) 435–447. [doi:10.1038/
s41586-025-08790-w](https://doi.org/10.1038/s41586-025-08790-w).
- [5] F. Melozzi, M. M. Woodman, V. K. Jirsa, C. Bernard, The Virtual
890 Mouse Brain: A computational neuroinformatics platform to study
whole mouse brain dynamics, *eNeuro* 4 (3) (2017) ENEURO.0111–
17.2017. [doi:10.1523/ENEURO.0111-17.2017](https://doi.org/10.1523/ENEURO.0111-17.2017).
- [6] K. Dai, S. L. Gratiy, Y. N. Billeh, R. Xu, B. Cai, N. Cain, A. E. Rime-
895 haug, A. J. Stasik, G. T. Einevoll, S. Mihalas, C. Koch, A. Arkhipov,
Brain Modeling ToolKit: an open source software suite for multiscale
modeling of brain circuits, *PLOS Computational Biology* 16 (11) (2020)

- e1008386. [doi:10.1371/journal.pcbi.1008386](https://doi.org/10.1371/journal.pcbi.1008386).
- [7] Y. N. Billeh, B. Cai, S. L. Gratiy, K. Dai, R. Iyer, N. W. Gouwens, et al.,
 900 Systematic integration of structural and functional data into multi-scale
 models of mouse primary visual cortex, *Neuron* 106 (3) (2020) 388–
 403.e18. [doi:10.1016/j.neuron.2020.01.040](https://doi.org/10.1016/j.neuron.2020.01.040).
- [8] H. Xiong, C. Chu, L. Fan, M. Song, J. Zhang, Y. Ma, R. Zheng, J. Zhang,
 Y. Zhengyi, T. Jiang, The Digital Twin Brain: A bridge between bio-
 905 logical and artificial intelligence, *Intelligent Computing* 2 (2023) 0055.
[doi:10.34133/icomputing.0055](https://doi.org/10.34133/icomputing.0055).
- [9] Z. Chen, C. Li, W. Lu, X. Xue, J. Feng, T. Zeng, MouseDTB: A mouse
 digital twin brain at single-neuron resolution, [OpenReview preprint](#),
 ICLR 2026 submission, modified February 2026 (2026).
- 910 [10] D. Wadden, S. Lin, K. Lo, L. L. Wang, M. van Zuylen, A. Cohan,
 H. Hajishirzi, Fact or fiction: Verifying scientific claims, in: *Proceedings
 of the 2020 Conference on Empirical Methods in Natural Language Pro-
 cessing*, Association for Computational Linguistics, 2020, pp. 7534–7550.
[doi:10.18653/v1/2020.emnlp-main.609](https://doi.org/10.18653/v1/2020.emnlp-main.609).
- 915 [11] D. Wadden, K. Lo, L. L. Wang, A. Cohan, I. Beltagy, H. Hajishirzi,
 MultiVerS: Improving scientific claim verification with weak supervision
 and full-document context, in: *Findings of the Association for Compu-
 tational Linguistics: NAACL 2022*, Association for Computational Lin-
 guistics, 2022, pp. 61–76. [doi:10.18653/v1/2022.findings-naacl.6](https://doi.org/10.18653/v1/2022.findings-naacl.6).
- 920 [12] X. Lu, L. Pan, Q. Liu, P. Nakov, M.-Y. Kan, SciTab: A challenging
 benchmark for compositional reasoning and claim verification on scien-
 tific tables, in: *Proceedings of the 2023 Conference on Empirical Meth-*

- ods in Natural Language Processing, Association for Computational Linguistics, 2023, pp. 7787–7813. [doi:10.18653/v1/2023.emnlp-main.483](https://doi.org/10.18653/v1/2023.emnlp-main.483).
- [13] C. Wang, Y. Shen, Z. Kuang, A. Cohan, Y. Zhao, SciVer: Evaluating foundation models for multimodal scientific claim verification, in: Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics, Association for Computational Linguistics, 2025, pp. 8562–8579. [doi:10.18653/v1/2025.acl-long.420](https://doi.org/10.18653/v1/2025.acl-long.420).
- [14] S. R. Javaji, Y. Cao, H. Li, Y. Yu, N. Muralidhar, Z. Zhu, Can AI validate science? benchmarking LLMs for accurate scientific claim–evidence reasoning, arXiv (2025). [arXiv:2506.08235](https://arxiv.org/abs/2506.08235), [doi:10.48550/arXiv.2506.08235](https://doi.org/10.48550/arXiv.2506.08235).
- [15] B. G. Buchanan, E. H. Shortliffe (Eds.), Rule-Based Expert Systems: The MYCIN Experiments of the Stanford Heuristic Programming Project, Addison-Wesley, Reading, MA, 1984.
- [16] R. Studer, V. R. Benjamins, D. Fensel, Knowledge engineering: Principles and methods, Data & Knowledge Engineering 25 (1–2) (1998) 161–197. [doi:10.1016/S0169-023X\(97\)00056-6](https://doi.org/10.1016/S0169-023X(97)00056-6).
- [17] M. Mohamad, J.-P. Steghöfer, R. Scandariato, Security assurance cases—state of the art of an emerging approach, Empirical Software Engineering 26 (2021) 70. [doi:10.1007/s10664-021-09971-7](https://doi.org/10.1007/s10664-021-09971-7).
- [18] T. Lebo, S. Sahoo, D. McGuinness, PROV-O: The PROV Ontology, World Wide Web Consortium, W3C Recommendation, 30 April 2013 (2013).
URL <https://www.w3.org/TR/prov-o/>

- [19] H. Knublauch, D. Kontokostas, [Shapes Constraint Language \(SHACL\)](#), World Wide Web Consortium, W3C Recommendation, 20 July 2017 (2017).
 950 URL <https://www.w3.org/TR/shacl/>
- [20] F. A. Batarseh, L. Freeman, C.-H. Huang, A survey on artificial intelligence assurance, *Journal of Big Data* 8 (2021) 60. doi:10.1186/s40537-021-00445-7.
- 955 [21] American Society of Mechanical Engineers, Assessing Credibility of Computational Modeling Through Verification and Validation: Application to Medical Devices, New York, NY, ASME V&V 40-2018 (2018).
- [22] J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d’Alché Buc, E. Fox, H. Larochelle, [Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program](#), *Journal of Machine Learning Research* 22 (164) (2021) 1–20.
 960 URL <https://www.jmlr.org/papers/v22/20-303.html>
- [23] A. Fernández-Isabel, I. Martín de Diego, R. Fuentes-Fernández, Mouse-BrainBench: A claim-aware validation framework for partial mouse-brain digital models, companion manuscript under review. Supporting software and artifacts available at <https://github.com/Fiutten/Mouse-brain> (2026).
 965
- [24] J. Vladika, F. Matthes, Scientific fact-checking: A survey of resources and approaches, in: Findings of the Association for Computational Linguistics: ACL 2023, Association for Computational Linguistics, 2023, pp. 6215–6230. doi:10.18653/v1/2023.findings-acl.387.
 970
- [25] D. Wadden, K. Lo, B. Kuehl, A. Cohan, I. Beltagy, L. L. Wang, H. Ha-

- jishirzi, SciFact-Open: Towards open-domain scientific claim verification, in: Findings of the Association for Computational Linguistics: EMNLP 2022, Association for Computational Linguistics, 2022, pp. 4719–4734. [doi:10.18653/v1/2022.findings-emnlp.347](https://doi.org/10.18653/v1/2022.findings-emnlp.347).
- [26] N. Tan, M. Li, M. Gahegan, SciTrue: Evidence-grounded claim verification in science, in: Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations, Association for Computational Linguistics, 2026, pp. 397–406. [doi:10.18653/v1/2026.eacl-demo.27](https://doi.org/10.18653/v1/2026.eacl-demo.27).
- [27] J. Chen, W. Zhan, H. Wang, B. Mayer, D. Fu, D. Zhou, PhyVer: Physics-grounded material claim verification with multi-fidelity physical evidence, in: Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics: System Demonstrations, Association for Computational Linguistics, 2026, pp. 546–555. [doi:10.18653/v1/2026.acl-demo.54](https://doi.org/10.18653/v1/2026.acl-demo.54).
- [28] D. C. Zhang, S. Cui, Z. Chu, X. Zhang, D. Lee, MEVER: Multi-modal and explainable claim verification with graph-based evidence retrieval, in: Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics, Association for Computational Linguistics, 2026, pp. 5236–5255. [doi:10.18653/v1/2026.eacl-long.242](https://doi.org/10.18653/v1/2026.eacl-long.242).
- [29] Y. Geifman, R. El-Yaniv, [SelectiveNet: A deep neural network with an integrated reject option](#), in: Proceedings of the 36th International Conference on Machine Learning, Vol. 97 of Proceedings of Machine Learning Research, PMLR, 2019, pp. 2151–2159.

URL <https://proceedings.mlr.press/v97/geifman19a.html>

- [30] I. Benbasat, J. S. Dhaliwal, A framework for the validation of knowledge acquisition, *Knowledge Acquisition* 1 (2) (1989) 215–233. [doi:10.1016/S1042-8143\(89\)80020-2](https://doi.org/10.1016/S1042-8143(89)80020-2).
1000
- [31] R. Knauf, A. J. Gonzalez, T. Abel, A framework for validation of rule-based systems, *IEEE Transactions on Systems, Man, and Cybernetics, Part B* 32 (3) (2002) 281–295. [doi:10.1109/TSMCB.2002.999805](https://doi.org/10.1109/TSMCB.2002.999805).
- [32] F. Ikhwantri, D. Marijan, Evaluating assurance cases as text-attributed graphs for structure and provenance analysis, *arXiv* (2026). [arXiv:2604.20577](https://arxiv.org/abs/2604.20577), [doi:10.48550/arXiv.2604.20577](https://doi.org/10.48550/arXiv.2604.20577).
1005
- [33] M. T. Moghaddam, Compliance-by-construction argument graphs: Using generative AI to produce evidence-linked formal arguments for certification-grade accountability, in: 2026 IEEE Workshop on Formal Arguments for CPS Certification (FACCT), IEEE, 2026, pp. 21–26.
1010 [doi:10.1109/FACCT71761.2026.00009](https://doi.org/10.1109/FACCT71761.2026.00009).
- [34] R. A. Rasheed, S. Banerjee, A. Mukherjee, R. Hazra, From fluent to verifiable: Claim-level auditability for deep research agents, *arXiv* (2026).
1015 [arXiv:2602.13855](https://arxiv.org/abs/2602.13855), [doi:10.48550/arXiv.2602.13855](https://doi.org/10.48550/arXiv.2602.13855).
- [35] National Academies of Sciences, Engineering, and Medicine, Reproducibility and Replicability in Science, The National Academies Press, Washington, DC, 2019. [doi:10.17226/25303](https://doi.org/10.17226/25303).
- [36] M. D. Wilkinson, M. Dumontier, I. J. Aalbersberg, G. Appleton, M. Axton, A. Baak, N. Blomberg, J.-W. Boiten, L. B. da Silva Santos, P. E. Bourne, et al., The FAIR guiding principles for scientific data management and stewardship, *Scientific Data* 3 (2016) 160018.
1020

[doi:10.1038/sdata.2016.18](https://doi.org/10.1038/sdata.2016.18).

- 1025 [37] M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, T. Gebru, Model cards for model reporting, in: Proceedings of the Conference on Fairness, Accountability, and Transparency, 2019, pp. 220–229. [doi:10.1145/3287560.3287596](https://doi.org/10.1145/3287560.3287596).
- [38] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, K. Crawford, Datasheets for datasets, Communications
1030 of the ACM 64 (12) (2021) 86–92. [doi:10.1145/3458723](https://doi.org/10.1145/3458723).
- [39] M. Barker, N. P. Chue Hong, D. S. Katz, A.-L. Lamprecht, C. Martinez-Ortiz, F. Psomopoulos, J. Harrow, L. J. Castro, M. Gruenpeter, P. A. Martinez, T. Honeyman, Introducing the fair principles for research software, Scientific Data 9 (2022) 622. [doi:10.1038/s41597-022-01710-x](https://doi.org/10.1038/s41597-022-01710-x).
- 1035 [40] O. Rübel, A. Tritt, R. Ly, B. K. Dichter, S. Ghosh, L. Niu, P. Baker, I. Soltesz, L. Ng, K. Svoboda, L. Frank, K. E. Bouchard, The Neurodata Without Borders ecosystem for neurophysiological data science, eLife 11 (2022) e78362. [doi:10.7554/eLife.78362](https://doi.org/10.7554/eLife.78362).
- [41] W. L. Oberkamp, C. J. Roy, Verification and Validation in Scientific Computing, Cambridge University Press, 2010. [doi:10.1017/
1040 CB09780511760396](https://doi.org/10.1017/CB09780511760396).
- [42] R. G. Sargent, An introduction to verification and validation of simulation models, in: Proceedings of the 2013 Winter Simulation Conference, 2013, pp. 321–327. [doi:10.1109/WSC.2013.6721430](https://doi.org/10.1109/WSC.2013.6721430).
- 1045 [43] S. Riedmaier, B. Danquah, B. Schick, F. Diermeyer, Unified framework and survey for model verification, validation and uncertainty quantification, Archives of Computational Methods in Engineering 28 (2021)

- 2655–2688. [doi:10.1007/s11831-020-09473-7](https://doi.org/10.1007/s11831-020-09473-7).
- [44] J. E. Knox, K. D. Harris, N. Graddis, J. D. Whitesell, H. Zeng, J. A. Harris, E. Shea-Brown, S. Mihalas, High-resolution data-driven model of the mouse connectome, *Network Neuroscience* 3 (1) (2019) 217–236. [doi:10.1162/netn_a_00066](https://doi.org/10.1162/netn_a_00066).
- [45] P. Turishcheva, P. G. Fahey, M. Vystrčilová, L. Hansel, R. Froebe, K. Ponder, Y. Qiu, K. F. Willeke, M. Bashiri, R. Baikulov, Y. Zhu, L. Ma, S. Yu, T. Huang, B. M. Li, W. De Wulf, N. Kudryashova, M. H. Hennig, N. L. Rochefort, A. Onken, E. Wang, Z. Ding, A. S. Tolias, F. H. Sinz, A. S. Ecker, Retrospective for the Dynamic Sensorium Competition for predicting large-scale mouse primary visual cortex activity from videos, *arXiv* (2024). [arXiv:2407.09100](https://arxiv.org/abs/2407.09100), [doi:10.48550/arXiv.2407.09100](https://doi.org/10.48550/arXiv.2407.09100).
- [46] E. Y. Wang, et al., Foundation model of neural activity predicts response to new stimulus types, *Nature* 640 (2025) 470–477. [doi:10.1038/s41586-025-08829-y](https://doi.org/10.1038/s41586-025-08829-y).
- [47] D. Liscai, E. Luconi, A. Marin Vargas, A. Sanzeni, Beyond single neurons: population response geometry in digital twins of mouse visual cortex, *Journal of Statistical Mechanics: Theory and Experiment* 2025 (9) (2025) 094003. [doi:10.1088/1742-5468/adde43](https://doi.org/10.1088/1742-5468/adde43).
- [48] Z. Ding, P. G. Fahey, S. Papadopoulos, et al., Functional connectomics reveals general wiring rule in mouse visual cortex, *Nature* 640 (2025) 459–469. [doi:10.1038/s41586-025-08840-3](https://doi.org/10.1038/s41586-025-08840-3).
- [49] J. M. Mooij, J. Peters, D. Janzing, J. Zscheischler, B. Schölkopf, [Distinguishing cause from effect using observational data: Methods and](#)

- benchmarks, Journal of Machine Learning Research 17 (32) (2016) 1–102.
- 1075 URL <https://jmlr.org/papers/v17/14-518.html>
- [50] M. R. Lynn, Determination and quantification of content validity, Nursing Research 35 (6) (1986) 382–385.
- [51] C. B. Terwee, C. A. C. Prinsen, A. Chiarotto, M. J. Westerman, D. L. Patrick, J. Alonso, L. M. Bouter, H. C. W. de Vet, L. B. Mokkink, 1080 COSMIN methodology for evaluating the content validity of patient-reported outcome measures: A delphi study, Quality of Life Research 27 (5) (2018) 1159–1170. doi:10.1007/s11136-018-1829-0.
- [52] R. Hasan, R. Sion, M. Winslett, [The case of the fake picasso: Preventing history forgery with secure provenance](#), in: 7th USENIX Conference on File and Storage Technologies, USENIX Association, San Francisco, CA, 2009. 1085
- URL <https://www.usenix.org/conference/fast09/technical-sessions/presentation/hasan>
- [53] S. E. Lazic, The problem of pseudoreplication in neuroscientific studies: Is it affecting your analysis?, BMC Neuroscience 11 (2010) 5. doi:10.1186/1471-2202-11-5. 1090
- [54] E. Neyhart, N. Zhou, B. R. Munn, R. G. Law, C. Smith, Z. H. Mridha, F. A. Blanco, G. Li, Y. Li, M. Hu, M. J. McGinley, J. M. Shine, J. Reimer, Cortical acetylcholine dynamics are predicted by cholinergic axon activity and behavior state, Cell Reports 43 (10) (2024) 114808. 1095
- doi:10.1016/j.celrep.2024.114808.
- [55] D. J. Millman, G. K. Ocker, S. Caldejon, I. Kato, J. D. Larkin, E. K.

Lee, J. Luviano, C. Nayan, T. V. Nguyen, K. North, S. Seid, C. White,
J. Lecoq, C. Reid, M. A. Buice, S. E. J. de Vries, VIP interneurons in
1100 mouse primary visual cortex selectively enhance responses to weak but
specific stimuli, eLife 9 (2020) e55130. [doi:10.7554/eLife.55130](https://doi.org/10.7554/eLife.55130).